

## Chapter 3: Classification

- Intro notes
  - MNIST data set: “Hello world!” of ml classification
  - Some ML algos sensitive to order and perform poorly if many similar instances in a row
  - Shuffling solves that. But shuffling is not good for time series data!
- Training a binary classifier
  - Stochastic Gradient Descent
- Performance Measures
  - Training a classifier trickier than a regressor
  - Sometimes need more than sklearn cross-validation → implement own validation
  - Don’t use Accuracy for skewed datasets
  - Measuring accuracy using cross-validation
  - Confusion matrix
    - Accuracy:  $(TP+TN)/(TP + FP + TN + FN)$
    - Precision:  $TP/(TP + FP)$
    - Recall:  $TP/(TP + FN)$ 
      - AKA: Sensitivity, True Positive Rate
  - Precision and recall
    - F-score: harmonic mean of precision & recall
      - Low scores get more weight
      - Only high F if both precision & recall high
      - $2 * \frac{precision * recall}{precision + recall}$
      - $\frac{TP}{TP + \left(\frac{FN + FP}{2}\right)}$
    - Predict bad videos to protect kids → favor high precision
    - Predict shoplifters → favor high recall
  - Precision/Recall Trade-off
    - If out all your positive picks, most were correct: high precision
    - If out of all positive instances, you picked most of them: high recall
    - If most of your positive picks were correct, but your positive picks did not capture many positive instances: high precision/low recall
    - Decision threshold: achieve high precision or recall
      - Sometimes threshold ↑, precision ↓ due to size of data set
      - Recall can only ↓ when threshold increases → higher requirement to classify as positive
    - Some asks for high precision, respond: “At what recall?”
  - The ROC Curve
    - True Positive vs. false positive rate
    - $FPR = 1 - TNR$ ; TNR AKA specificity
    - Sensitivity (TPR) vs.  $1 - \text{Specificity (TNR)}$
    - Compare classifiers: Area Under the Curve
      - Perfect = 1, Random = 0.5

- PR or ROC?
    - PR if +ve class is rare or care more about false positives than false negatives
- Multiclass Classification
  - Algos: SGC, Random Forest, Naïve Bayes; Logistic and SVM are binary
  - Ways to train multiclass using binary:
    - One-versus-rest (or all): train as many classifiers as labels, pick the one with highest decision score
      - Preferred for most binary algos
    - One-versus-one: train one classifier for every pair; N classes:  $N*(N-1)/2$  classifiers
      - OvO preferred for large datasets with algos that don't scale well
- Error Analysis
  - SGD classifier is linear model: assigns weight per class to each data point
    - Seeing a new piece of info, sums up weighted intensities to get a score
- Multilabel classification
  - Distinguishing multiple objects within one space
  - Evaluate using f1 score (or other binary classifier metric) for each label, then average
    - Assumes all labels equally important
- Multioutput classification
  - Generalization of multilabel where each label can be multiclass